

## MODEL OPTIMIZATION

# Bias, Variance, and Overfitting Explained, Step by Step

You have likely heard about bias and variance before. They are two fundamental terms in machine learning and often used to explain overfitting and underfitting. If you're working with machine learning methods, it's crucial to understand these concepts well so that you can make optimal decisions in your own projects. In this article, you'll learn everything you need to know about bias, variance, overfitting, and the bias-variance tradeoff.



BORIS GIBA

17 APR 2021 • 20 MIN READ

Last modified: 17 Apr 2021



Background image by Sora Shimazaki ([link](#))

# Outline

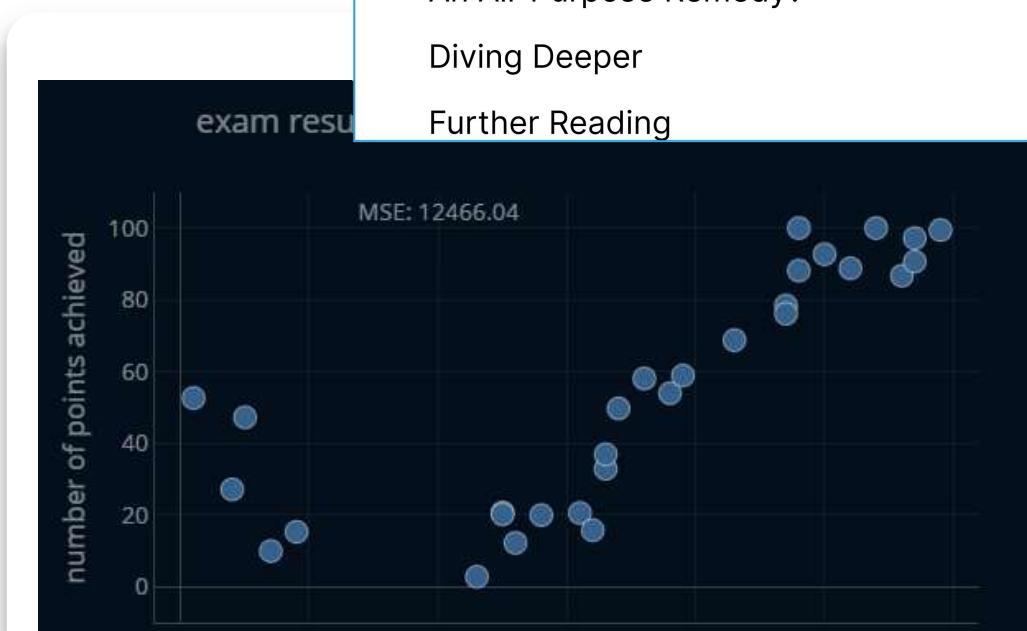
Bias and variance are very fundamental, and also very important concepts. Understanding bias and variance well will help you make more effective and more well-reasoned decisions in your own machine learning projects, whether you're working on your personal portfolio or at a large organization. In this article, you will learn what bias and variance are, what the so-called *bias-variance tradeoff* is, and how you can make machine learning projects better.

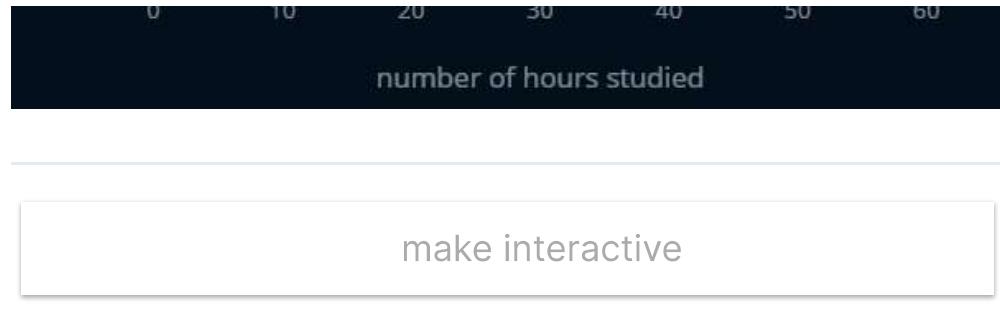
## The Problem

Let's start by looking at a simple problem. You're teaching your friend for an exam. You give them a small dataset to practice with. The dataset contains the number of hours spent studying for the exam (between 0 and 100) and the corresponding points achieved (between 0 and 100). You then tell your friend to predict the points achieved based on the number of hours spent studying. The dataset looks like this:

### Table of Contents

- Outline
- The Problem
- The Bias
  - Practical Examples (Bias)
  - Revisiting the Definition of Bias
  - But how do we tell if our model is good?
- The Variance
  - Practical Examples (Variance)
  - Revisiting the Definition of Variance
- The Tradeoff
  - The Tradeoff?
  - An All-Purpose Remedy?
- Diving Deeper
- Further Reading





make interactive

Your friend is just getting started with machine learning and recently read [an article about linear regression](#). Now they want to try out linear regression themselves! After some time, they come back to you creating a model that fits the data well, but the predicted values alone don't tell the whole story. By adding the number of hours studied to the equation, it makes the results easier to digest, the model more interpretable. The resulting function. The resulting function.

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

make interactive

You look at the plot and notice a couple of things.

1. There definitely seems to be a relationship between the number of hours studied and the number of points achieved.
2. This relationship does seem somewhat linear, especially when considering the segment between 40 and 60 hours studied.
3. However, the linear relationship is not that apparent throughout the entire graph. When you take a look at the interval between 0 and 30 hours studied, it does seem to have more of a ~~downward trend in contrast to~~ the interval between 40 and 60 hours studied.

With linear regression, we can fit a linear function to model the relationship between our features (number of hours studied) and the target variable (number of points achieved). But what if we want to capture the same trend throughout the entire range of the data, but we don't want to throw off one half of the data? In this case, we need to use a non-linear regression model called polynomial regression. If we can perform a simple linear regression instead. We can see that there are other relationships in our data than just linear ones. You can always fit a linear model, even if they go ahead and tell you that it's not good.

Imagine you are in the situation where you have to find out the best degree for your polynomial model, meaning the maximum power which you will apply to your features. Is a degree of 2 or 3 already enough? Do you need a degree of 10 or higher? Pause for a second and think about this question. Below you can find an interactive visualization where you can take a look at all the possible polynomial regression plots for the degrees 1-20. After you've come up with a number in your head, take a look at the visualization and see if your intuition was right.\*

## Table of Contents

- Outline
- The Problem
- The Bias
  - Practical Examples (Bias)
  - Revisiting the Definition of Bias
  - But how do we tell if our model is good?
- The Variance
  - Practical Examples (Variance)
  - Revisiting the Definition of Variance
- The Tradeoff
  - The Tradeoff?
- An All-Purpose Remedy?
- Diving Deeper
- Further Reading

If you estimated the congrats! You eyes, can tell that the best fits the data. But the degree of our polynomials always gets closer a model with degree 2 than the middle than the degree is, the “wiggly” model with a higher can take on more complicated function shapes. So maybe the model with degree=20 is truly the best one? Ok, let’s take a step back.

We have all of these different functions and some can capture the relationship between our features and our target better than others. If we were to compare the models with degree=1 and degree=20 respectively, we would say that the one with degree=20 “captures the relationship better”. But in a

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

practical scenario, there has to be a better way of expressing this without having to say “captures this relationship better” every time. Wouldn’t it be more satisfying if we could somehow say that the model with degree=20 has a higher/lower *something* than the model with degree=1? We can do exactly that by introducing the notion of **bias**.

## The Bias

We want the bias to express how well a certain machine learning model fits a ~~particular dataset. So let's try and come~~ up with a formal definition. Then, we examples. Then, we in statistics and where from statistics. Our c

*The bias of a specific learning model can be defined as the difference between the expected value of the model's predictions and the true values of the target variable. In other words, it measures how far off our model's predictions are from the actual data points.*

That seems like a reasonable definition.

## Practical Examples

Let's take a look at three ways of fitting them using our new polynomial regression models with degrees 1, 4, and 15 respectively. Below is the graph showing our dataset and the predictions of the model with a degree of 1.

## Table of Contents

- Outline
- The Problem
- The Bias
  - Practical Examples (Bias)
  - Revisiting the Definition of Bias
  - But how do we tell if our model is good?
- The Variance
  - Practical Examples (Variance)
  - Revisiting the Definition of Variance
- The Tradeoff
  - The Tradeoff?
  - An All-Purpose Remedy?
  - Diving Deeper
  - Further Reading

As you see, this model's predicted values are therefore we can say because it does not fit the number of points studied.

In practice, we probably displaying our prediction bias is based on how we want to compare two concrete way of measuring it we often use the (train) make deductions about a based task in front of us squared error to measure the performance of a specific model. To make the results a bit easier to interpret, let's also add a square root on top so we can immediately tell by how much the average point is approximately off. If you are wondering why we don't use a different metric like the mean absolute error or why we don't use the root mean squared error by default, all of this is explained in more detail in this segment of the article about linear regression.

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

If we compute the root mean squared error (RMSE) for the predictions of this model, we get a value of 19.8. This means that one single point can be off by up to ~19.8, i.e. if a student in truth achieved 80 points, our model might only give them 60, which would probably not make the student very happy. With this we can capture the following behavior:

- large training error → large bias

Let's now look at the model with degree 4:

The image shows a table of contents for a machine learning course. The main title is "Table of Contents". Below it is a list of topics, each with a small icon to its left. The topics are: Outline, The Problem, The Bias (with subtopics Practical Examples (Bias) and Revisiting the Definition of Bias), The Variance (with subtopics Practical Examples (Variance) and Revisiting the Definition of Variance), The Tradeoff (with subtopics The Tradeoff? and An All-Purpose Remedy?), Diving Deeper, and Further Reading.

	<b>Table of Contents</b>
	Outline
	The Problem
	The Bias
	Practical Examples (Bias)
	Revisiting the Definition of Bias
	But how do we tell if our model is good?
	The Variance
	Practical Examples (Variance)
	Revisiting the Definition of Variance
	The Tradeoff
	The Tradeoff?
	An All-Purpose Remedy?
	Diving Deeper
	Further Reading

This model predicts our  $y$  a lot better than the first one. It still makes some errors, but it does its job a lot better than the first model. If we compute the RMSE for the predictions of this model, we get 7.68. It's about 2.5 times better than the previous model! If a student in truth achieved 80 points, our model might only give them 72 (or up to 88), which still is not ideal, but it's considerably better than the previous model. We achieved this decrease in error by just moderately increasing

our degree. We've only increased the degree from 1 to 4, but managed to reduce our error by a factor of around 2.5! So this increase in complexity was certainly worth it. Because our model has a rather small error, we can say that it has a small bias since it does its task relatively well. With this we can capture the following behavior:

- small training error → small bias

Let's now take a closer look at the model with degree 15:

## Table of Contents

- Outline
- The Problem
- The Bias
  - Practical Examples (Bias)
  - Revisiting the Definition of Bias
  - But how do we tell if our model is good?
- The Variance
  - Practical Examples (Variance)
  - Revisiting the Definition of Variance
- The Tradeoff
  - The Tradeoff?
  - An All-Purpose Remedy?
  - Diving Deeper
- Further Reading

This model does look a bit weirder than the previous two, but it does in fact predict our data even better than the second one. If we compute the RMSE for the predictions of this model, we get 4.93. It's about 1.5 times lower than the previous error. However, we did increase our degree from 4 to 15, which is quite a lot. So this increase in complexity certainly is not as good of a deal as the previous one. But a

lower is a lower error, right? So technically, this model is the best of the bunch.

Because our model has a very low error, we can say that it has a very low bias since it does its task very well. With this we can capture the following behavior:

- very small training error → very small bias

## Revisiting the Definition of Bias

Ok, so we have definitely looked at a couple other ways that error is correlated with bias. The lower the error, the lower the bias. That's what we mean by bias.

When talking about bias, we usually talk about one model. We've always talked about one specific model and one particular metric.

Let's now look at the definition of bias to see if our previous understanding was correct. [Wikipedia](#):

*Statistical bias is the difference between the expected value of the results when all possible samples are taken that could be drawn from the same population, and a single observed sample result. In statistics, this concept is often used to characterize the properties of an estimator in terms of its sample statistics versus the true values it is estimating.*

## Table of Contents

- Outline
- The Problem
- The Bias
  - Practical Examples (Bias)
  - Revisiting the Definition of Bias
  - But how do we tell if our model is good?
- The Variance
  - Practical Examples (Variance)
  - Revisiting the Definition of Variance
- The Tradeoff
  - The Tradeoff?
  - An All-Purpose Remedy?
  - Diving Deeper
  - Further Reading

If we look back on how we introduced the bias, this makes sense. We compared the predictions of our model with the actual values present in our dataset. But the real definition is worded just a tiny bit differently, and contains two very important words:

Statistical bias is a feature of a statistical technique or of its results whereby the expected value of the results differs from the **true underlying** quantitative parameter being estimated.

This means that the bias is a way of describing the difference between the actual, *true* relationship in our data, and the one our model learned. In our examples, we've looked at the error between our predictions and the data points. Sure, that is a very sensible way to measure the bias of our machine learning models. But accurately representing features and the target contain noise, meaning the true relationship of data points is not always clear on it. Since I've created a linear regression function\*, I know the relationship between the features and the target. Below, I'll show you how to find a particular function. I'll also show you the button with which you can generate a new dataset.

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

make interactive

If the dataset contains no noise at all, it means that the dataset represents the *exact* underlying relationship. But almost every dataset out there will have some noise in it. This might be because the dataset was not properly cleaned, or maybe because there is some intrinsic noise to the underlying problem itself. Often times it is simply not possible to get perfect, noiseless data. For this reason, there might be an *irreducible error*, an error, that no machine learning model (or even human) can un-

statistics compares i  
underlying relationshi  
Because oftentimes,

I can show you how  
the noise because I  
mathematical function  
sampled from actual  
shape of some math  
it's exactly the oppos  
usually have a seemi  
create a mathematic  
relationship in that d  
created a function an  
using that function. I  
would have no way o  
what part of the data

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

This is the reason why we compare our dataset to our predictions. Ideally, we want our function to learn this true, underlying relationship, by only looking at the noisy dataset. Almost always we won't be able to create a model that perfectly matches this relationship (and we also don't have a way of checking if it did), but oftentimes we can generate a function that is *good enough*, assuming the dataset itself is *good enough*. So what does this mean? If our model creates a

relationship that we think is not as sensible, 4/5 times it will be the dataset that is to blame, and not the model.

## But how do we tell if our model is good?

Naturally, we are interested in keeping the bias as low as possible, because we want our model to make the best predictions possible. But if you take a look at the second and third model, would you really choose the third model over the second one? It does have a lower bias, so this is certainly an argument. But there is just something off about the third model. Maybe it's that it's too good at its task. There has to be a way that's not so great, or where to do this, we are not being careful with the variance.

## The Variance

With the bias, we now move on to the variance of machine learning models. The variance is the difference between feature and target values. We always only consider the variance of the training set, so we bring in values outside of the training set. This happens when we bring in new data that we won't consider a datapoint. For example, a model was trained to predict the exam score of a student from a dataset that is very similar to our current one. This could be the data of a different exam and different students from the same school. Alternatively, we can split our existing dataset into a *training*-set and a *testing*-set. We won't go into detail about why this makes sense in this particular article, but if you want to read more about this, I'd recommend you take a look at the article [Training and Testing Datasets Explained, Step by Step](#).

## Table of Contents

- Outline
- The Problem
- The Bias
  - Practical Examples (Bias)
  - Revisiting the Definition of Bias
- But how do we tell if our model is good?
- The Variance
  - Practical Examples (Variance)
  - Revisiting the Definition of Variance
- The Tradeoff
  - The Tradeoff?
- An All-Purpose Remedy?
- Diving Deeper
- Further Reading

We want the variance to express how *consistent* a certain machine learning model is in its predictions when compared across *similar* datasets. So let's apply the same approach that we used for the bias and try to first come up with a formal definition of variance ourselves. It might look like the definition below.

*The variance of a specific machine learning model trained on a specific dataset describes how much the performance of the machine learning model differs when evaluated at different points in time or on different datasets of similar origin. So for our model, the variance would tell us how much our particular model would change the value of a prediction if we evaluated it on a different dataset. In other words, the variance of our model would tell us how consistent our performance is across different datasets.*

Ok, now let's take a look at how we can implement regression models that have different degrees of variance.

## Practical Examples

This time, I'll add a function to my linear regression code that will calculate the variance of the dataset when compared against the information. First, we'll calculate the variance of the model and the (potentially) overfitted model. We'll also track the relative error of the model. We'll compare the MSE and the initial MSE of the model with a degree=1:

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

This model predicts is also maintained if the button a couple and the R.DIFF, you'll much. To be more exact difference of 1000 R you will receive a value change our dataset, which makes the model performs poorly makes bad predictions the performance of the multiple similar datasets variance.

With this we can capture

- small fluctuations

make interactive

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

In practice, we usually don't want to randomly alter our dataset and keep track of the relative difference to find out whether our variance is low or high. So how can we estimate the variance of a machine learning model? We look at the difference in performance between the training-set and the testing-set (or for that matter, the training set and any other dataset of similar origin). If this difference is high, so is the variance. If it is low, so is the variance.

Because the model with degree=1 has a high bias but a low variance, we say that it is *underfitting*, meaning it is not “fit enough” to accurately model the relationship between features and targets.

Let's now look at the model with degree 4:

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

This model predicts good performance is dataset. If you again keep an eye on the RMSE and the R.DIFF, you'll notice that the values change a little bit, but not a lot. To be more exact, if you compute the average relative difference of 1000 RMSE values for slightly altered datasets, you will receive a value of ~13.6. This means that if we slightly change our dataset, our MSE will fluctuate by around 13.6%, which makes the model *somewhat* consistent. It's certainly not as consistent as the first model, but it's not terrible as well. In other terms, our

model performs well, and does so pretty consistently. It always makes at least decent predictions for a dataset of this origin. Because the performance of this model is somewhat consistent across multiple similar datasets, we can say that it has a somewhat low variance.

With this we can capture the following behavior:

- medium fluctuation of the error → medium variance

Because the model has low variance, we say that there is a good balance between bias and variance.

Let's now look at the case where

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

make interactive

This model predicts our data very well. However, this outstanding performance is not at all maintained when we change up our dataset a bit. If you press the button just a few

times and look at the RMSE and R.DIFF, chances are you'll see some larger numbers than before. These numbers also change quite a lot between clicks. To be more exact, if you compute the average relative difference of 1000 RMSE values for slightly altered datasets, you will receive a value of ~169.5. This means that if we slightly change our dataset, our RMSE will fluctuate by around **169.5%**, which makes the predictions of this model as reliable as your office or home printer in the moment you need it most. This means that our model performs well sometimes, and catastrophically other times. Because inconsistent across runs, it has high variance.

With this we can capture:

- high fluctuation of error

Because this model is able to fit the training data perfectly, we say that it is *overfitted*. This means that it is able to fit this very exact dataset, but it fails to capture the underlying relationship that is true in the population.

## Revisiting the Definition of Variance

Ok, so we have defined bias and variance. Now we know that the error is the difference between the predicted value and the actual value. The lower the error for a particular model, the better it is. The higher the error fluctuation, the worse it is.

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

When talking about the variance of a particular model, we always talk about one model, but *multiple* datasets. In practice, you would compare the model error on the training dataset and the error on the testing (or validation) dataset.

Let's now also look at the definition used in statistics. In the words of [Wikipedia](#):

*In probability theory and statistics, variance is the expectation of the squared deviation of a random variable from its mean. In other words, it measures how far a set of numbers is spread out from their average value.*

The important part is "*spread out from their average value*".

This means that the variance is a way of describing the difference between the expected (or average) value and the predicted value of our model. This definition is similar to the one we described earlier.

Naturally, we are interested in minimizing variance as much as possible, because we want our predictions across all data points to be as close to each other as possible. We can see that the third model is not the best one, but it also has a lower variance than the second model. The first model is the best model out of the three, because it is minimizing both bias and variance. The first model (degree=1) has the lowest variance, while the third model minimizes the bias.

## The Tradeoff

Let's think about how we can find a good way to achieve a low variance without overfitting our model or to train our model for a shorter amount of time. If we want to extract *more insights* from our dataset, we want our model to *learn more* from our data. The more our model learns from our data, the better it will be at solving its task.

And how might we decrease variance? The easiest way to decrease variance would be to pick a more simple model or to train our existing model for a shorter amount of time. In other words, we want to extract *fewer insights* from our dataset, we

## Table of Contents

- Outline
- The Problem
- The Bias
  - Practical Examples (Bias)
  - Revisiting the Definition of Bias
  - But how do we tell if our model is good?
- The Variance
  - Practical Examples (Variance)
  - Revisiting the Definition of Variance
- The Tradeoff
  - The Tradeoff?
  - An All-Purpose Remedy?
  - Diving Deeper
  - Further Reading

want our model to *learn less* from our data. The less our model learns from our data, the more general it will be, because it has only learned so much to base its predictions on.

A good analogy would be to think of a student preparing for their mathematics and physics exams. If the student does not prepare at all, they will do poorly in both exams (high bias). If they put a lot of time into their preparation, they will probably do pretty well in both exams (low bias, low variance).

However, if the student over-prepares, they will do worse again (overfitting). They will train themselves not just for the exams, but to be good at solving problems they have solved over and over again on the exam that they have prepared without any error. However, if they have not practiced, they will struggle answering it (low bias again).

So there seems to be a tradeoff between bias and variance. You can only reduce one at a time. As variance starts to increase, the graph like this one:

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

make interactive

## The Tradeoff?

If you have a rather simple model like polynomial regression, there certainly is a tradeoff between bias and variance. As you saw, increasing decreased bias consistently increased variance. After that step, the error starts to increase dramatically and the model becomes meaningless.

However, if you have a more complex model like a deep neural network, things are not so clear-cut. Oftentimes it seems like there is no tradeoff at all. You can get into too much detail and end up with a model that is good enough to fill a separate book. This is where you turn to some external resources like the book *Deep Learning* by Yoshua Bengio or the paper [A Modern History of Neural Networks](#), which explains the history of neural networks. If you are looking for a more detailed explanation, you might want to read the paper [A Modern History of Neural Networks](#), which was written by Yoshua Bengio.

wrote the aforementioned blog post. The paper covers the topic a bit more in-depth.

## Table of Contents

Outline

The Problem

The Bias

Practical Examples (Bias)

Revisiting the Definition of Bias

But how do we tell if our model is good?

The Variance

Practical Examples (Variance)

Revisiting the Definition of Variance

The Tradeoff

The Tradeoff?

An All-Purpose Remedy?

Diving Deeper

Further Reading

## An All-Purpose Remedy?

No matter what machine learning model you are training and what problem you are trying to solve, there is one thing you can do that will almost always result in a better, more stable machine learning model. What is it? Simply said, “getting

more (and better) data". There have been numerous papers written and studies performed\* about whether increasing data capacity (and quality) yields a greater positive effect on model performance than trying to come up with a more elaborate machine learning algorithm. While it's impossible to generalize the results to each and every application, the general notion seems to be that increasing data capacity yields greater results than choosing a more complicated model.

So what does this mean for you? If you've already tried out several different models and alternative models don't seem to improve model performance, then the issue is frequently poor data or a bad model.

## Diving Deeper

In this article, I briefly mentioned the concept of *irreducible error*, an error that is inherent in itself, and that no matter what you do can undergo. This is something that might be amazing in theory, but in practice, it makes sense. Therefore it makes sense to start by analyzing the bias of your model. Models have been shown to stop improving any further no matter what you do, it may just be the case that you've reached the ideal bias for this particular problem and that it is simply not possible to improve any further with this particular dataset that you have.

Oftentimes you won't know the irreducible error for a problem. It is pretty difficult to compute and it's even more difficult to prove that it really is irreducible. In practice, you

## Table of Contents

- Outline
- The Problem
- The Bias
  - Practical Examples (Bias)
  - Revisiting the Definition of Bias
  - But how do we tell if our model is good?
- The Variance
  - Practical Examples (Variance)
  - Revisiting the Definition of Variance
- The Tradeoff
  - The Tradeoff?
  - An All-Purpose Remedy?
  - Diving Deeper
  - Further Reading

can't compute the irreducible error for every problem you try to solve with machine learning. What you can do instead, is to train multiple different machine learning models\*, calculate the bias for each of those models and compare them to each other. If you have multiple models that are all optimized for that particular problem and they all possess very similar biases, you might not get any better results by optimizing your models.

## Further Reading

Bias and variance are two metrics that describe overfitting and underfitting. To understand what exactly overfitting and underfitting are, and how to actually prevent it in your machine learning models, read the post [How To Prevent Overfitting](#), where you will learn how to make your model more accurate without overfitting or as underfitting.

You can only improve your model's performance until it becomes very accurate. However, there are limits to how much you can improve your model's performance when you optimize it. Models are trained on data, so you can improve your data quality, but once you can do so, you should move on to [Optimization](#), where you will learn exactly how you can do so.

## Table of Contents

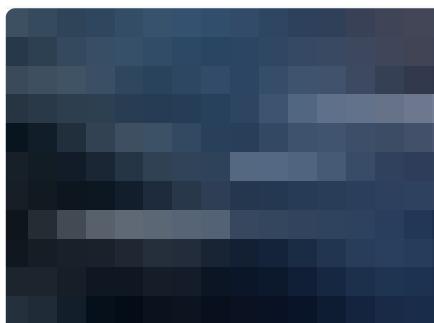
- Outline
- The Problem
- The Bias
  - Practical Examples (Bias)
  - Revisiting the Definition of Bias
  - But how do we tell if our model is good?
- The Variance
  - Practical Examples (Variance)
  - Revisiting the Definition of Variance
- The Tradeoff
  - The Tradeoff?
  - An All-Purpose Remedy?
  - Diving Deeper
- Further Reading



[0 comments](#)[Log In](#)

Comments powered by [Talkyard](#).

These posts may interest you:



#### MACHINE LEARNING MATH

## Gradient Descent for Linear Regression Explained, Step by Step

Gradient descent is one of the most famous techniques in machine learning and used for...



BORIS GIBA

03 APR 2021 • 23 MIN READ

## Table of Contents

[Outline](#)

[The Problem](#)

[The Bias](#)

[Practical Examples \(Bias\)](#)

[Revisiting the Definition of Bias](#)

[But how do we tell if our model is good?](#)

[The Variance](#)

[Practical Examples \(Variance\)](#)

[Revisiting the Definition of Variance](#)

[The Tradeoff](#)

[The Tradeoff?](#)

[An All-Purpose Remedy?](#)

[Diving Deeper](#)

[Further Reading](#)



BORIS GIBA

01 MAY 2021 • 15 MIN READ

[RSS](#)

[Reset Cookie Preferences](#)

[Imprint & Privacy Policy \(EN\)](#)

[Impressum & Datenschutz \(DE\)](#)

## Table of Contents

Outline

The Problem

The Bias

    Practical Examples (Bias)

    Revisiting the Definition of Bias

    But how do we tell if our model is good?

The Variance

    Practical Examples (Variance)

    Revisiting the Definition of Variance

The Tradeoff

    The Tradeoff?

    An All-Purpose Remedy?

    Diving Deeper

    Further Reading